

Business Problem - Predict the Price of Bangalore House

Using Linear Regression - Supervised Machine Learning Algorithm

Load Libraries

```
In [1]:  import pandas as pd
```

Load Data


```
In [2]:  path = r"https://drive.google.com/uc?export=download&id=1xxDtrZKfuWQf1-6K/
df = pd.read_csv(path)
```

```
In [3]:  df.head()
```

Out[3]:

	bath	balcony	price	total_sqft_int	bhk	price_per_sqft	area_typeSuper built-up Area	area_typeBuilt- up Area
0	3.0	2.0	150.0	1672.0	3	8971.291866	1	0
1	3.0	3.0	149.0	1750.0	3	8514.285714	0	1
2	3.0	2.0	150.0	1750.0	3	8571.428571	1	0
3	2.0	2.0	40.0	1250.0	2	3200.000000	1	0
4	2.0	2.0	83.0	1200.0	2	6916.666667	0	0

5 rows × 108 columns



Split Data

```
In [4]:  X = df.drop('price', axis=1)
y = df['price']

print('Shape of X = ', X.shape)
print('Shape of y = ', y.shape)
```

Shape of X = (7120, 107)
Shape of y = (7120,)

```
In [5]: ▶ from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, r

print('Shape of X_train = ', X_train.shape)
print('Shape of y_train = ', y_train.shape)
print('Shape of X_test = ', X_test.shape)
print('Shape of y_test = ', y_test.shape)

Shape of X_train = (5696, 107)
Shape of y_train = (5696,)
Shape of X_test = (1424, 107)
Shape of y_test = (1424,)
```

Feature Scaling

```
In [6]: ▶ from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
sc.fit(X_train)
X_train = sc.transform(X_train)
X_test = sc.transform(X_test)
```

Linear Regression - ML Model Training

```
In [7]: ▶ from sklearn.linear_model import LinearRegression
lr = LinearRegression()

lr.fit(X_train, y_train)
```

Out[7]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)

In [8]: lr.coef_

```

Out[8]: array([-5.70206143e+00, -1.25679916e+00,  8.27341833e+01, -1.44906911e+0
1,
           5.75662723e+01,  1.88468905e-01, -1.72593897e+00, -4.51058311e+0
0,
          -2.22589244e+00, -4.28978455e+00, -2.44590976e+00,  5.40246226e-0
1,
          -1.03633400e+00,  1.43064873e+00, -6.25029424e-02, -1.51548783e+0
0,
          -2.14422789e-01,  2.16244155e+00, -1.48710228e+00,  1.95250816e+0
0,
          -3.10761125e+00, -1.28138668e+00, -1.01367155e+00,  1.37968545e-0
2,
           1.10383858e+00,  1.26497611e+00, -3.52405517e+00, -1.21398741e+0
0,
          -5.04622019e-01,  1.46299181e+00, -5.50064233e-01, -8.46468162e-0
2,
           6.84882188e-01, -1.39849820e+00, -1.94761710e-02, -1.57716300e+0
0,
           4.20886278e-01,  8.03443207e-01,  2.99182164e+00,  3.86430413e-0
3,
           1.05037261e-01,  2.89115612e-01, -3.16916626e-01,  1.05625868e+0
0,
          -1.39649279e+00, -3.10533604e+00,  1.01764011e-01, -7.49672917e-0
2,
          -8.03271555e-01, -1.27061856e+00, -8.54046164e-01,  2.64566484e-0
1,
           9.10688839e-01, -8.23059458e-01, -9.07215234e-01,  1.22059216e+0
0,
           2.11418894e+00, -5.38187400e-01, -1.32164338e+00, -8.28349340e-0
1,
           1.28167980e+00, -1.92911295e-01,  6.65824485e-02,  3.65563139e-0
2,
          -1.85069853e+00,  1.49068024e+00, -9.57964753e-01, -9.36110163e-0
1,
          -7.45634897e-01,  7.22643165e-02, -6.79260144e-01, -1.70853833e-0
1,
          -1.72288643e+00, -1.15833746e+00,  5.78931788e-01,  1.37836966e+0
0,
          -1.14424496e+00,  3.96188294e-01, -6.08013157e-01, -2.20959218e+0
0,
           3.45270810e-01,  1.01747431e-03,  1.06563895e-01,  3.04728530e+0
0,
           2.09496392e+00, -8.13481923e-01, -4.18437282e-01,  2.30993396e+0
0,
           3.31858800e-02,  8.07865914e-02,  5.37064987e-02,  1.55347699e+0
0,
           8.13889657e-01, -1.14636462e+00,  3.41805788e-01, -8.28022037e-0
1,
           1.68897360e+00,  2.97657524e-01,  9.59437517e-01,  4.57297702e-0
1,
          -2.22729515e-01, -1.48290835e+00, -6.26342867e-01,  5.86538254e-0
1,
          -1.78547310e+00,  2.19020231e-01, -3.45032599e-01])

```

```
In [9]: lr.intercept_
```

```
Out[9]: 95.0802729985955
```

Predict the value of Home and Test

```
In [10]: X_test[0, :]
```

```
Out[10]: array([ 0.71301986,  0.0112734 ,  0.30202307,  0.65677518, -0.48064341,
                -1.7385623 ,  2.11587407, -0.25430867,  0.51007548, -0.18373025,
                -0.16389438, -0.1473229 , -0.13023539, -0.12812824, -0.12598816,
                -0.12454231, -0.12953656, -0.12381344, -0.12010681, -0.11551113,
                -0.10992018, -0.10909925, -0.10660036, -0.11234866, -0.09315135,
                -0.08618799, -0.08923672, -0.09023078, -0.08721571, -0.09023078,
                -0.08721571, -0.08195215, -0.08195215, -0.07633675, -0.0751646 ,
                -0.08085949, -0.0739743 , -0.07975227, -0.07153563, -0.0751646 ,
                -0.0677166 , -0.08085949, -0.07153563, -0.07862985, -0.0751646 ,
                -0.07862985, -0.06504853, -0.0751646 , -0.06901264, -0.0751646 ,
                -0.06901264, -0.07028523, -0.07276497, -0.07028523, -0.06367332,
                -0.06226825, -0.06226825, -0.06639573, -0.06504853, -0.05935999,
                -0.06083125, -0.06639573, -0.06639573, -0.06226825, -0.06367332,
                -0.05935999, -0.06639573, -0.06367332, -0.06226825, -0.06226825,
                -0.05935999, -0.05935999, -0.05935999, -0.05630391, -0.05935999,
                -0.05785186, -0.05935999, -0.05935999, -0.06083125, -0.06083125,
                -0.05471275, -0.06083125, -0.06226825, -0.05935999, -0.05935999,
                -0.06226825, -0.06226825, -0.05785186, -0.06504853, -0.06226825,
                -0.06083125, -0.05935999, -0.05307449, -0.05630391, -0.06226825,
                -0.05471275, -0.05935999, -0.05471275, -0.05471275, -0.05138463,
                -0.05307449, -0.05307449, -0.05471275, -0.05471275, -0.05630391,
                -0.05630391, -0.05138463])
```

```
In [11]: lr.predict([X_test[0, :]])
```

```
Out[11]: array([76.90661876])
```

```
In [12]: lr.predict(X_test)
```

```
Out[12]: array([ 76.90661876,  15.25005377, 113.6828165 , ...,  21.30296864,
                71.43462962, 230.0414626 ])
```

```
In [13]: ▶ y_test
```

```
Out[13]: 2435      80.00
          3113      40.00
          426     120.00
          1124     79.00
          1161     45.00
          ...
          2078     28.34
          6855     84.00
          4381     32.00
          3862     63.00
           43    180.00
          Name: price, Length: 1424, dtype: float64
```

```
In [14]: ▶ lr.score(X_test, y_test)
```

```
Out[14]: 0.7903837092682249
```

Implementing Ridge and Lasso Regression

```
In [15]: ▶ from sklearn.linear_model import Ridge, Lasso
```

```
In [16]: ▶ rd = Ridge()
          rd.fit(X_train,y_train)
          rd.score(X_test, y_test)
```

```
Out[16]: 0.790568637433663
```

```
In [18]: ▶ ls = Lasso()
          ls.fit(X_train,y_train)
          ls.score(X_test, y_test)
```

```
Out[18]: 0.8036373003525774
```

```
In [19]: ▶ rd2 = Ridge(alpha = 2)
          rd2.fit(X_train,y_train)
          rd2.score(X_test, y_test)
```

```
Out[19]: 0.7907530260397942
```

```
In [20]: ▶ ls2 = Lasso(alpha=2)

ls2.fit(X_train,y_train)

ls2.score(X_test, y_test)
```

Out[20]: 0.8160181533703601

```
In [21]: ▶ ls3 = Lasso(alpha=3)

ls3.fit(X_train,y_train)

ls3.score(X_test, y_test)
```

Out[21]: 0.8263450613017426

```
In [ ]: ▶
```